GETTING STARTED WITH DIGITS™

Allison Gray

Solutions Architect

AGENDA

- INTRODUCTION TO DIGITS AND IMAGE CLASSIFICATION
- HOW TO INSTALL DIGITS
- CREATING DATASETS
- TRAIN A NETWORK
- MODIFY YOUR NETWORK

NVIDIA DIGITS

Interactive Deep Learning GPU Training System

Process Data

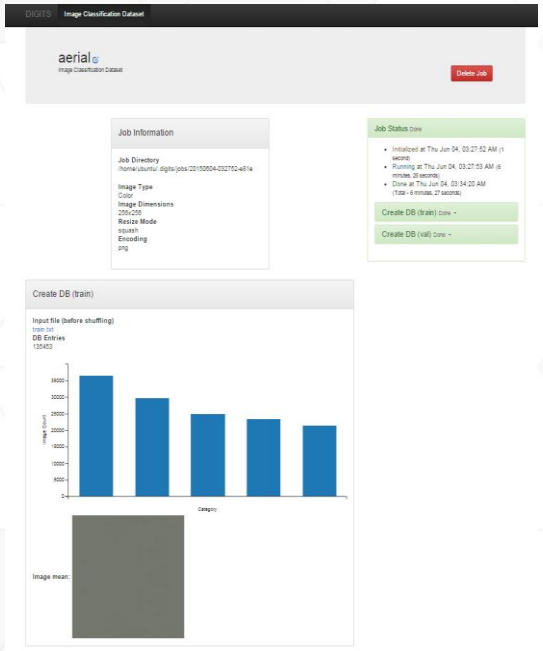

Configure DNN

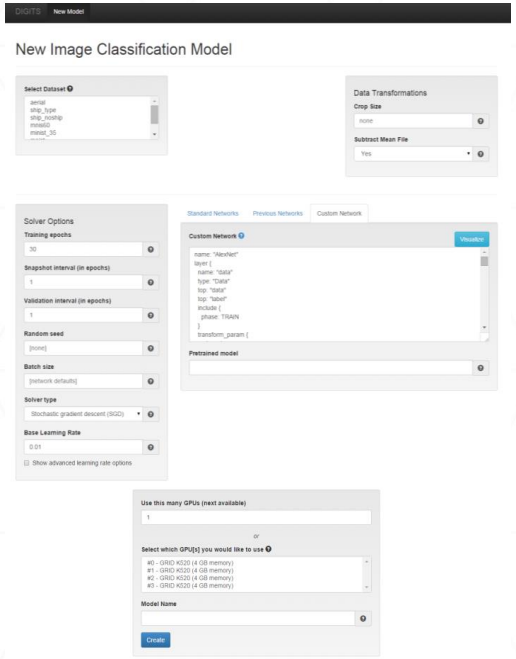

Monitor Progress

Visualization

NVIDIA DIGITS

Who is DIGITS for?

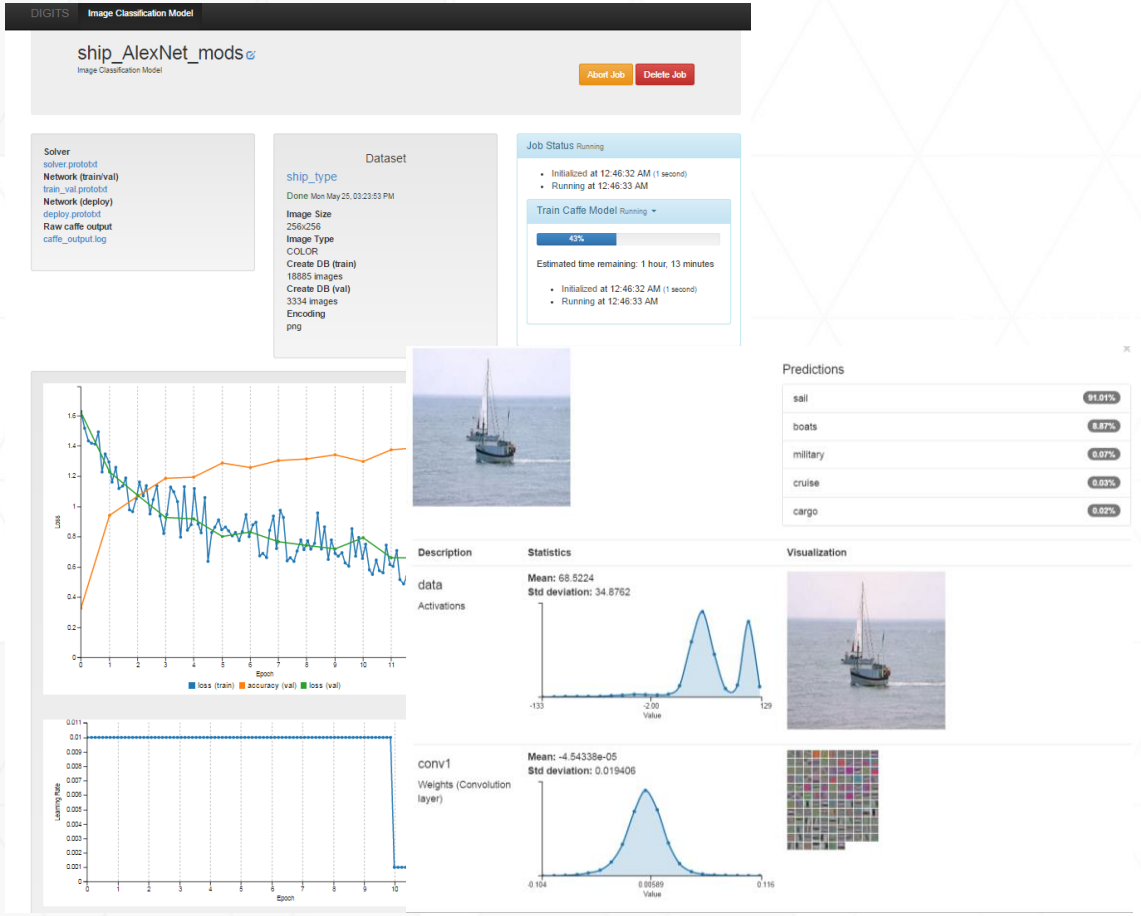

Data Scientists & Researchers:

- Quickly design the best deep neural network (DNN) for your data
- Monitor DNN training quality in real-time
- Manage training of many DNNs in parallel on multi-GPU systems, and multi-GPU training

WHAT IS DEEP LEARNING?

A Classification Example

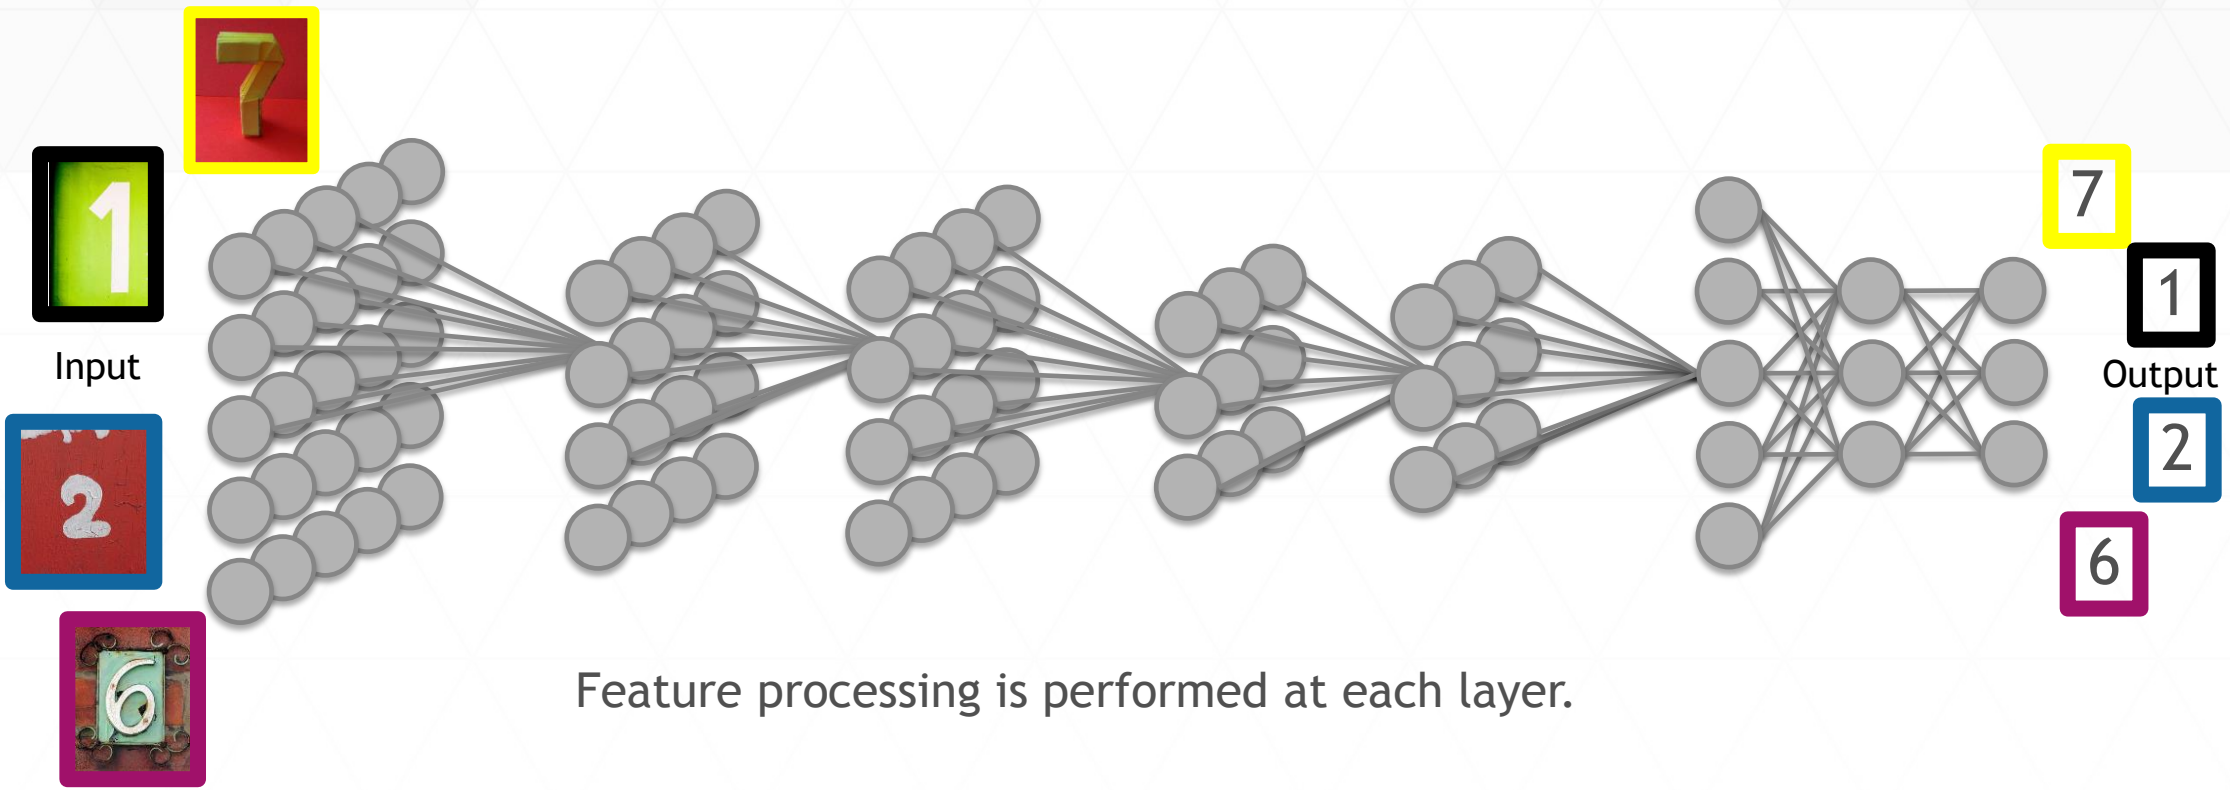

Feature processing is performed at each layer.

IMAGE CLASSIFICATION WITH DNNS

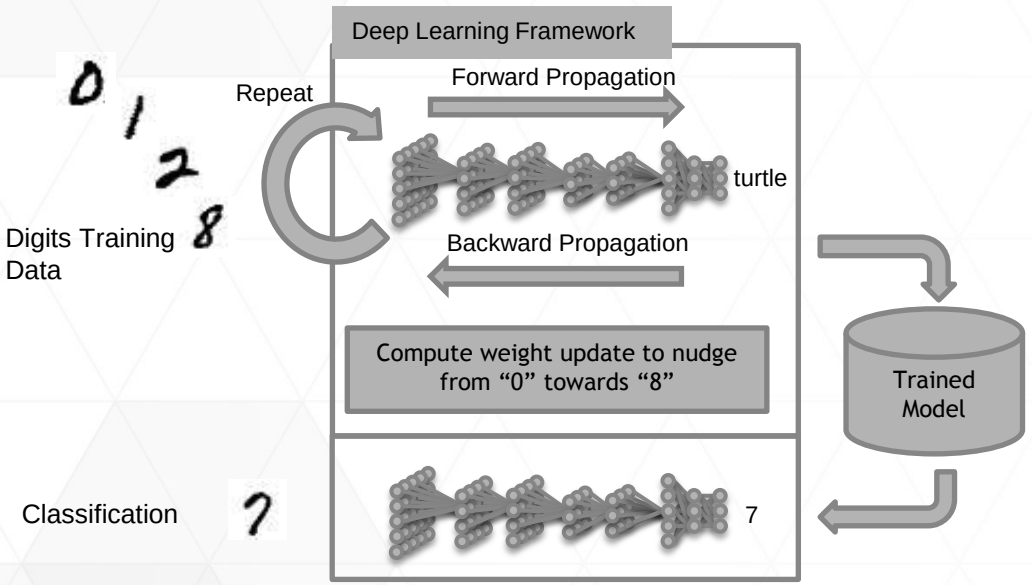

- Example training process
  - Pick a DNN design
  - Input training images
- Test accuracy
  - If bad: modify DNN, fix training set or update training parameters

HOW DO I GET DIGITS?

Two ways to install it

- Easy way
  - OS - Ubuntu 14.04
  - Download the web installer - https://developer.nvidia.com/digits
- Not as easy way
  - OSX, Windows(untested)
  - Build NVIDIA Caffe branch - https://github.com/NVIDIA/caffe
  - Download DIGITS from github – https://github.com/NVIDIA/DIGITS
- Hardware/software recommendations
  - GPU(s) with compute capabilities >=3.0 for cuDNN support
  - OS - Ubuntu 14.04

NVIDIA DIGITS

Hands-on Lab

Quickly train and improve a DNN well suited for your data

- Create a database
- Start out with a standard network and start training
  - Classify images to further understand performance
- Modify your network
- Adding/removing parameters
- Inflating your data

NVIDIA DIGITS

Creating your Dataset

Main Console

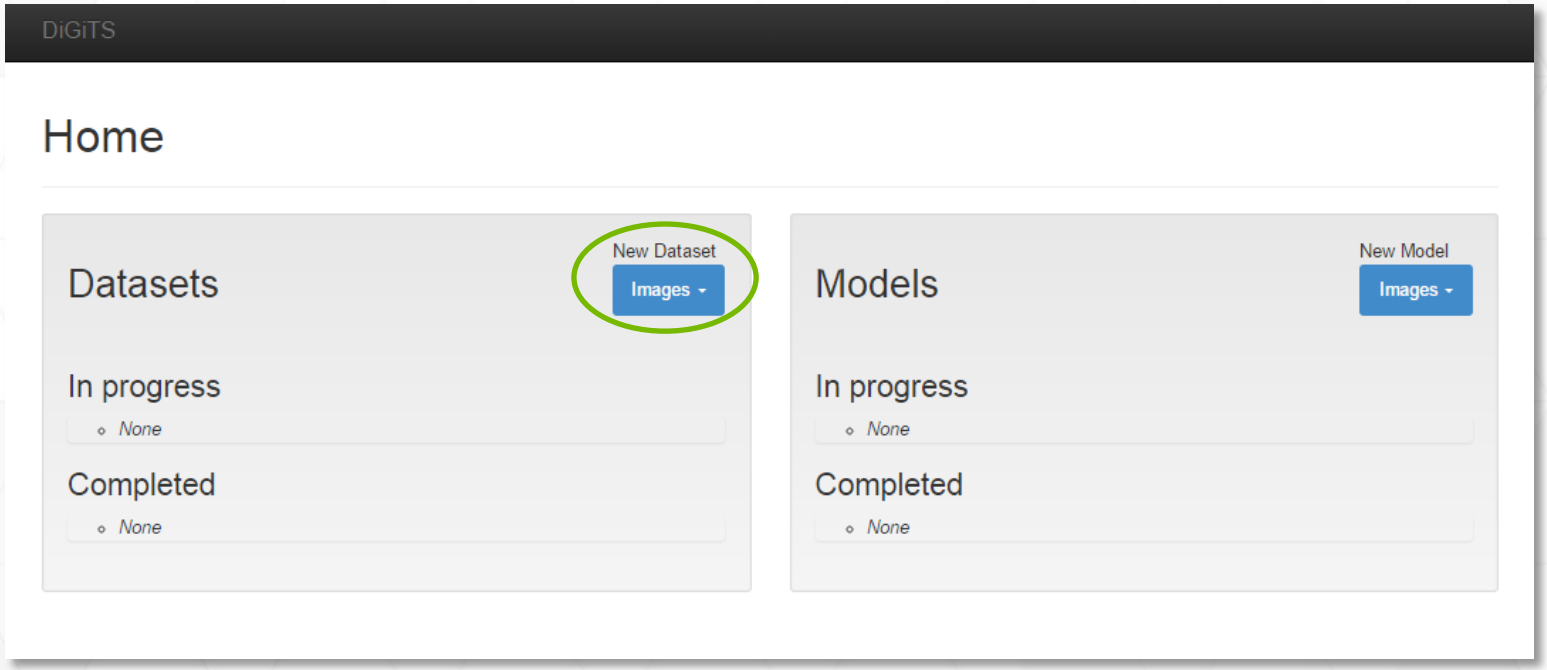

NVIDIA DIGITS

Create Database

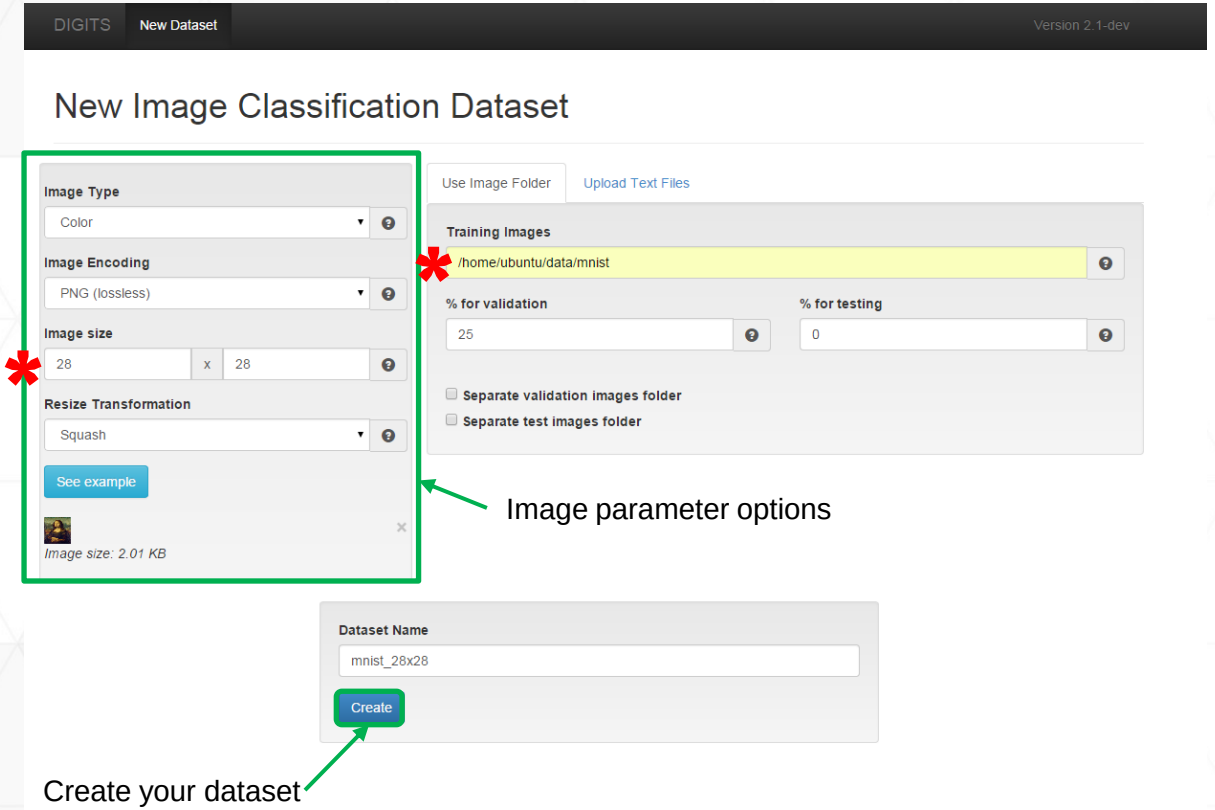

Input Data Format

```
/path/to/images/
                |-0/
                |    |-0_0jpg
                |    |-0_1.jpg
                |-1/
                |    |1_0.jpg
                |    |1_1.jpg
                |-2/
                     |-2_0.jpg
                     |-2_1.jpg
```

* Lab tasks

NVIDIA DIGITS

Database results

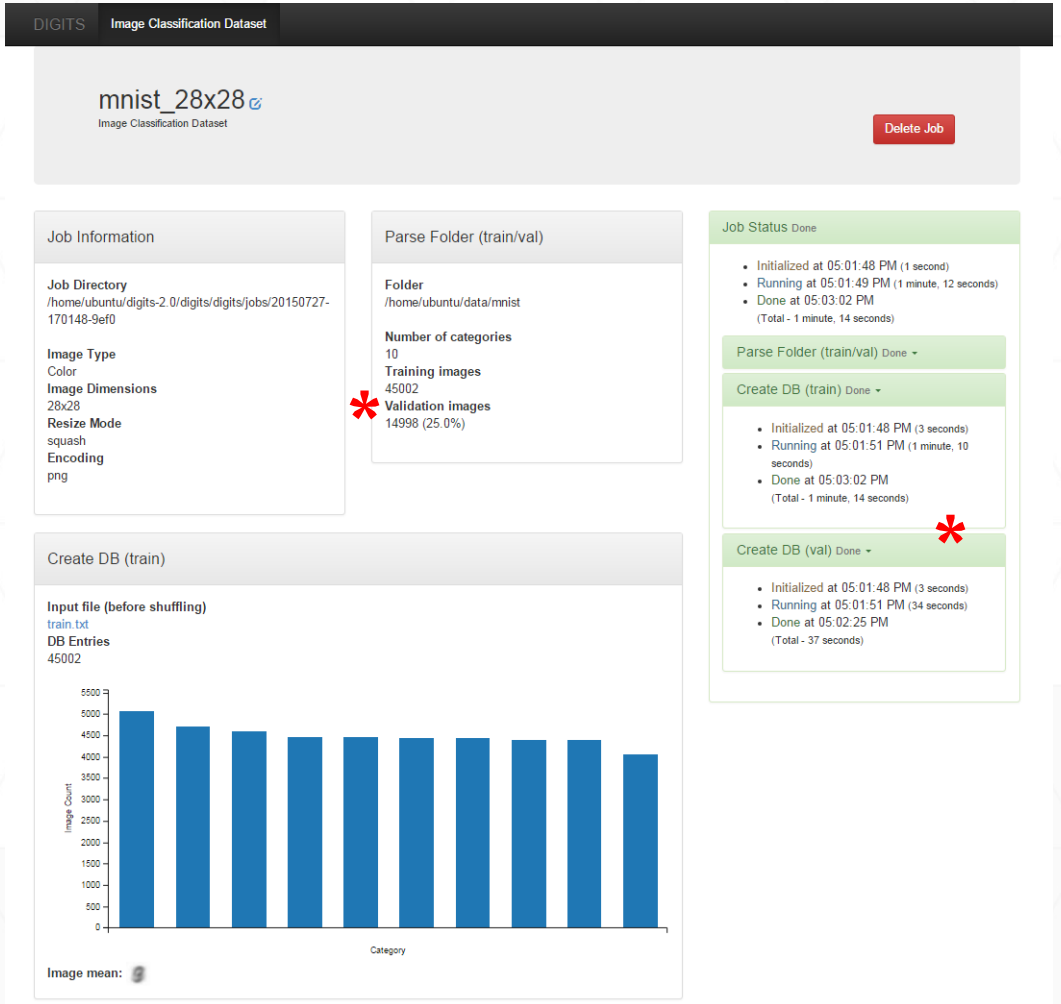

* Lab tasks

TRAINING AND VALIDATION DATA

Why do we need it?

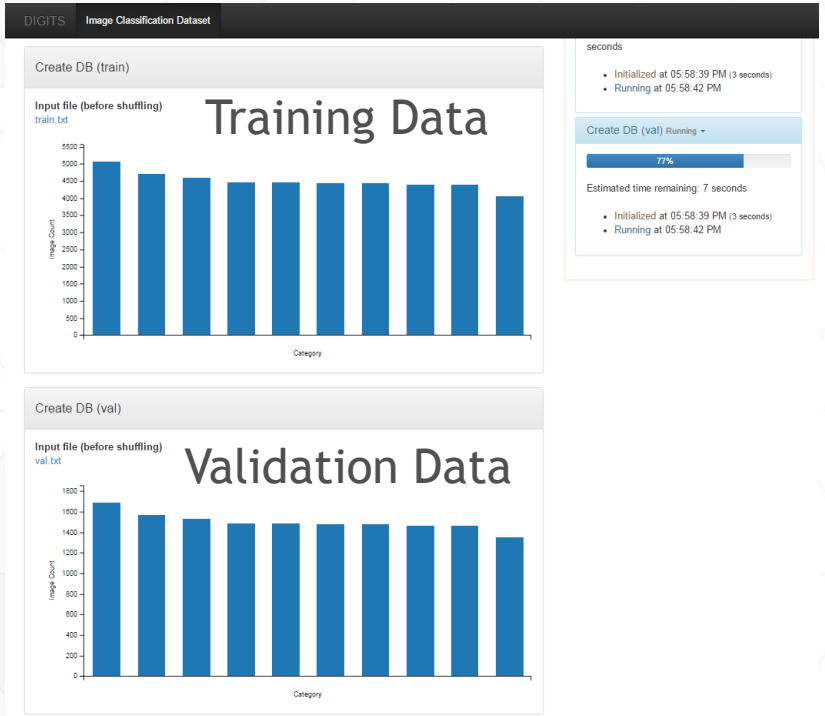

- Training data is used to train our neural network
  - Teaches the network to classify object categories
- Validation data tests the performance of the network
  - This data is only used for testing the generalization ability of the network
  - Not used to teach/train network
  - Prevents use of and identifies when network is overfitting

NVIDIA DIGITS

Network Configuration – Start with a Standard Network

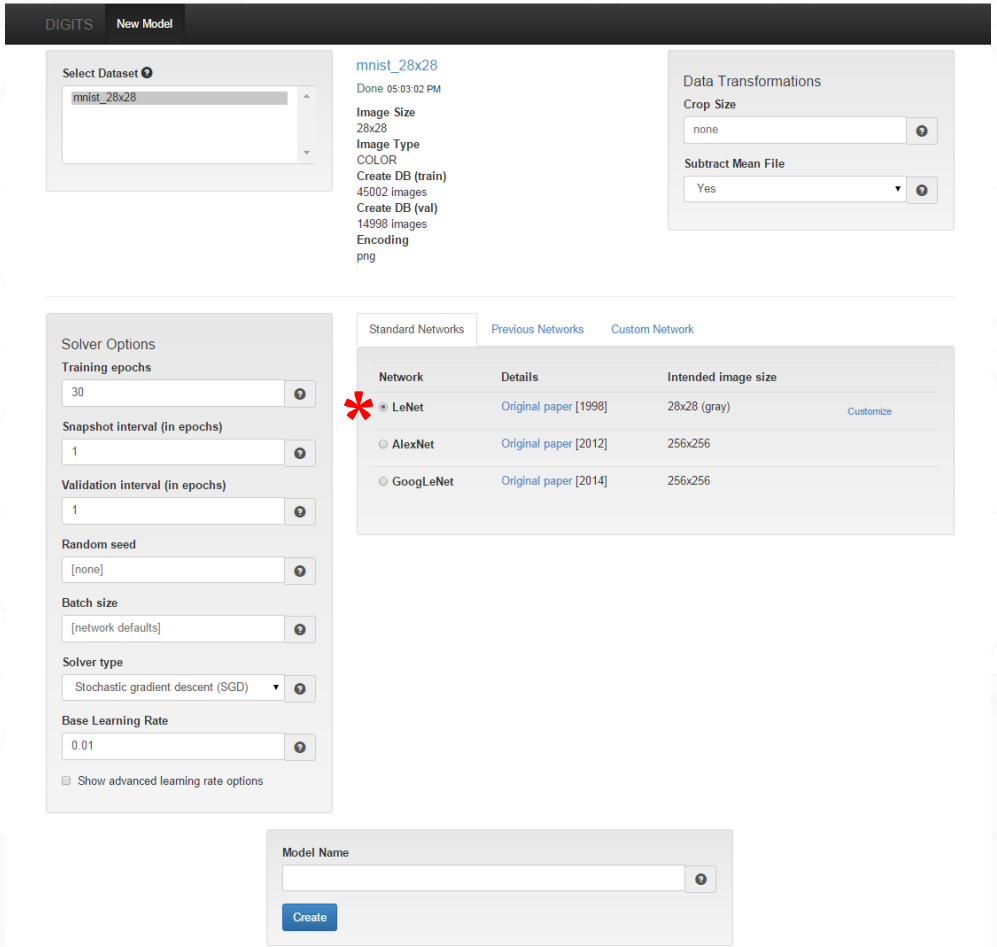

* Lab tasks

LENET

Network Configuration

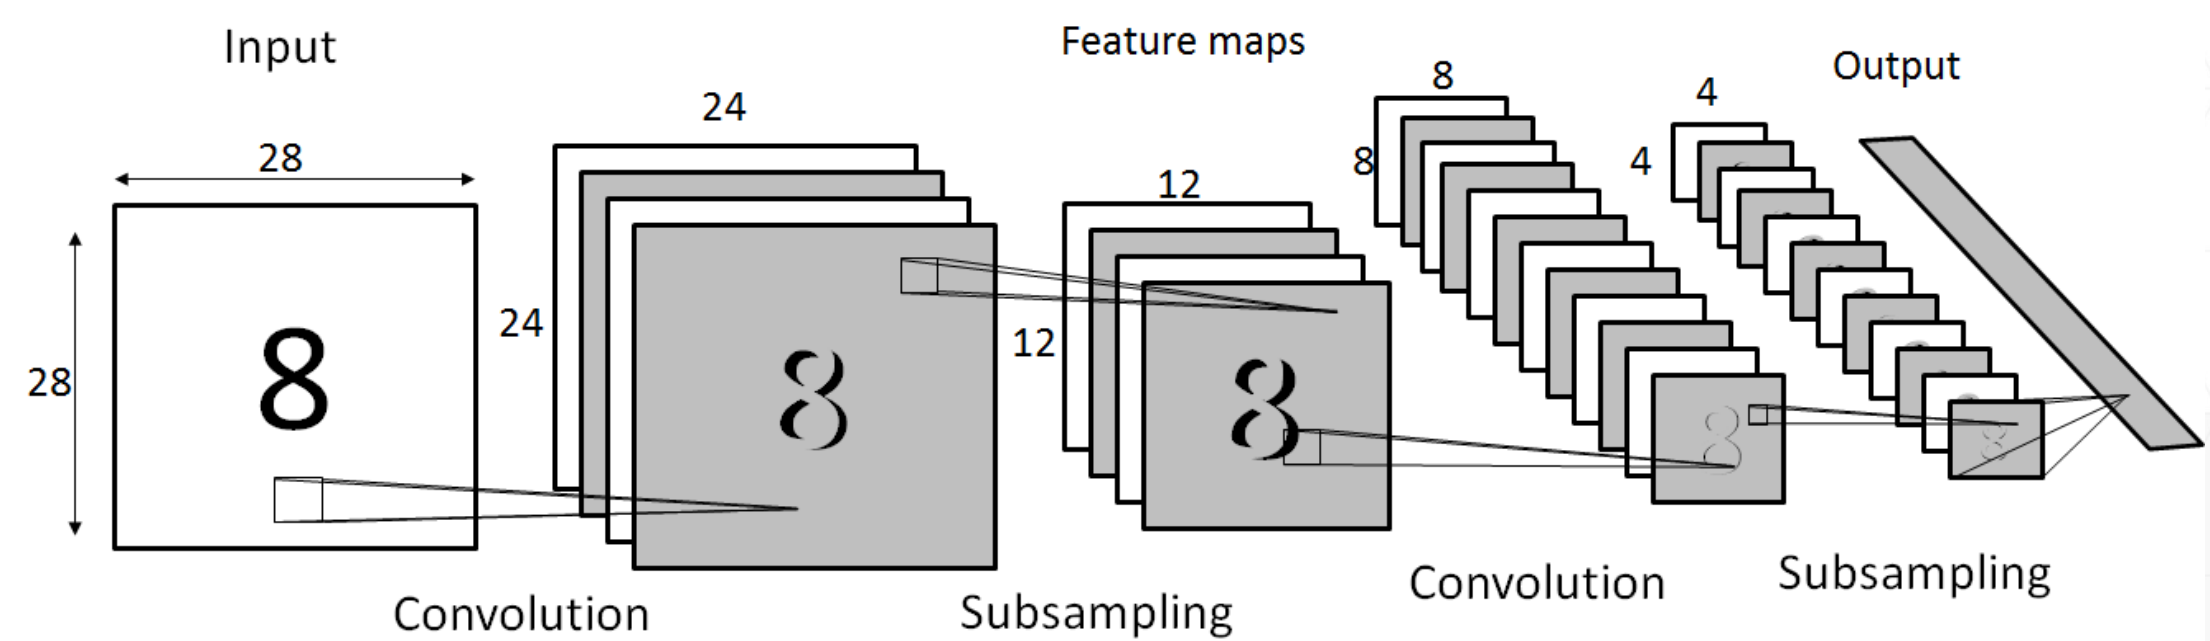

NVIDIA DIGITS

Visualizing your Network

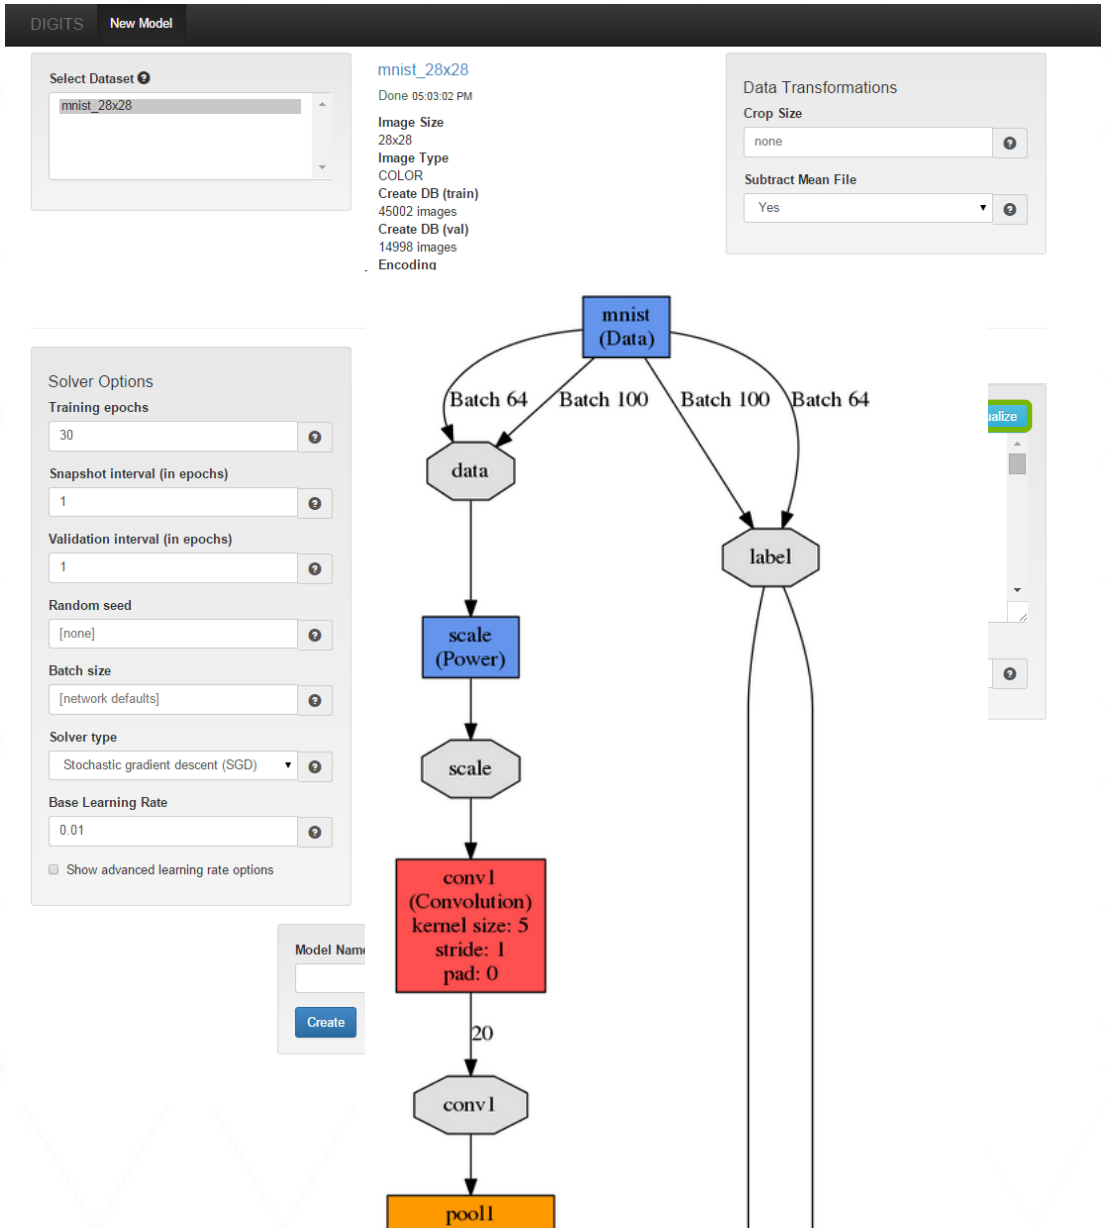

NETWORK PARAMETERS

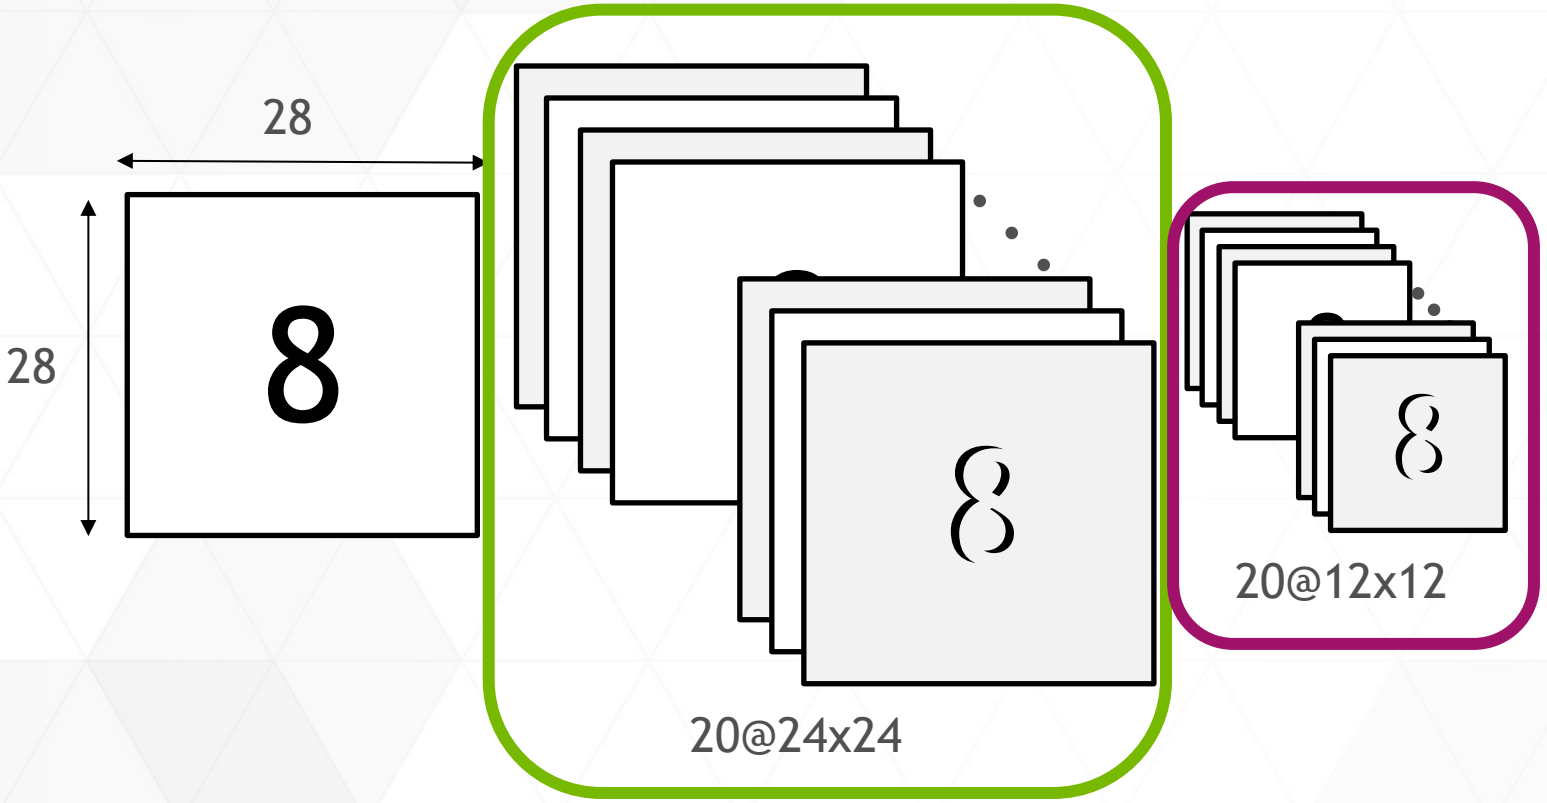

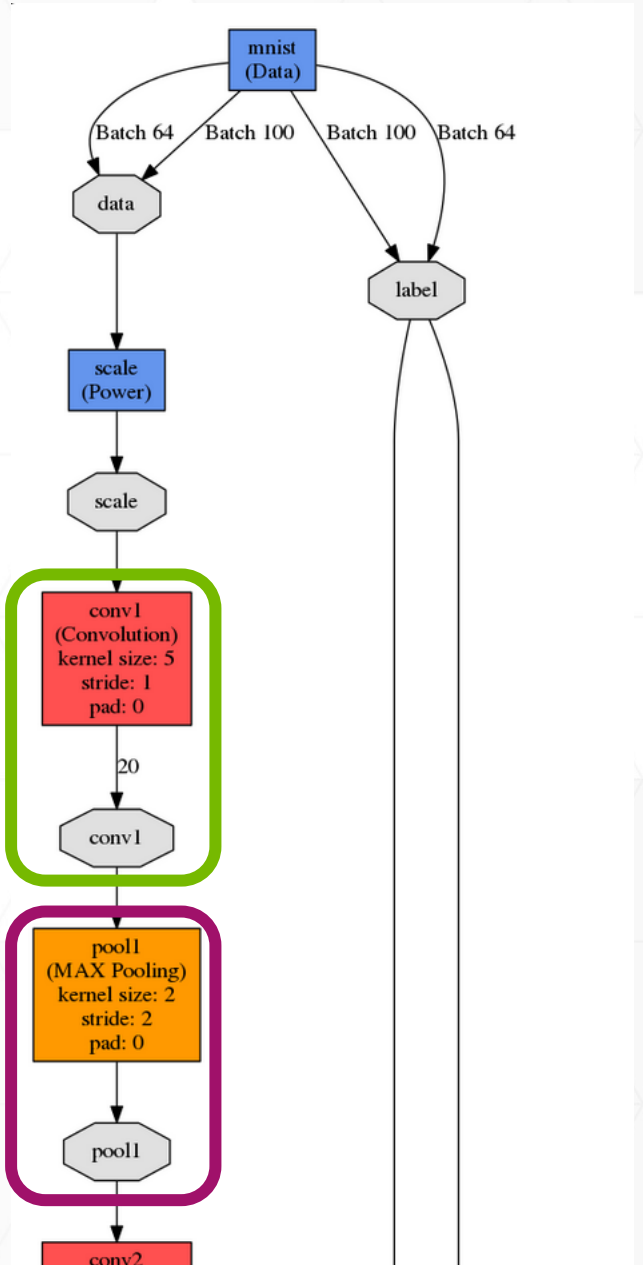

NVIDIA DIGITS

Solver Parameters

- Training epochs – processing of all data
- Snapshot interval – saving trained network
- Validation interval – DNN test with the validation data
- Batch size – number of images processed together
- Solver type – SGD, ADAGRAD,NAG
- Learning rate and policy

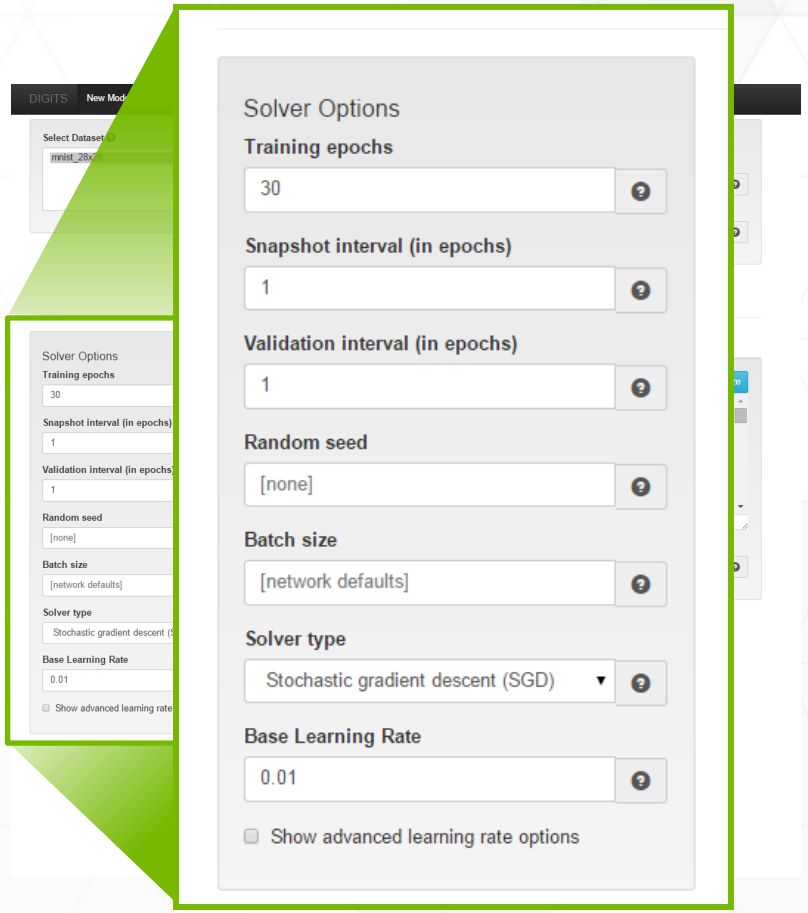

NVIDIA DIGITS

Single and multi-GPU training is easy

Singe GPU system

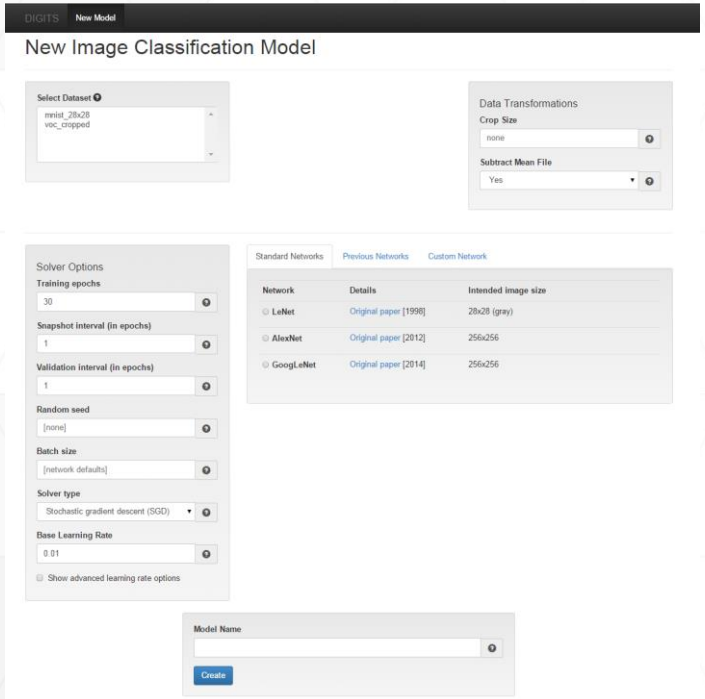

Multiple GPU system

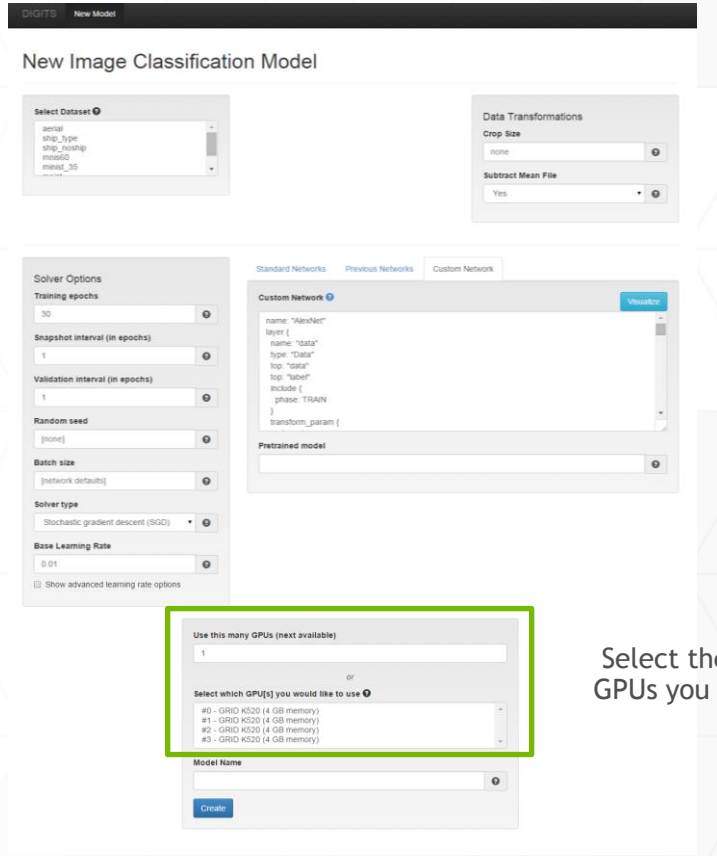

Select the number of GPUs you want to use.

NVIDIA DIGITS

Training Results

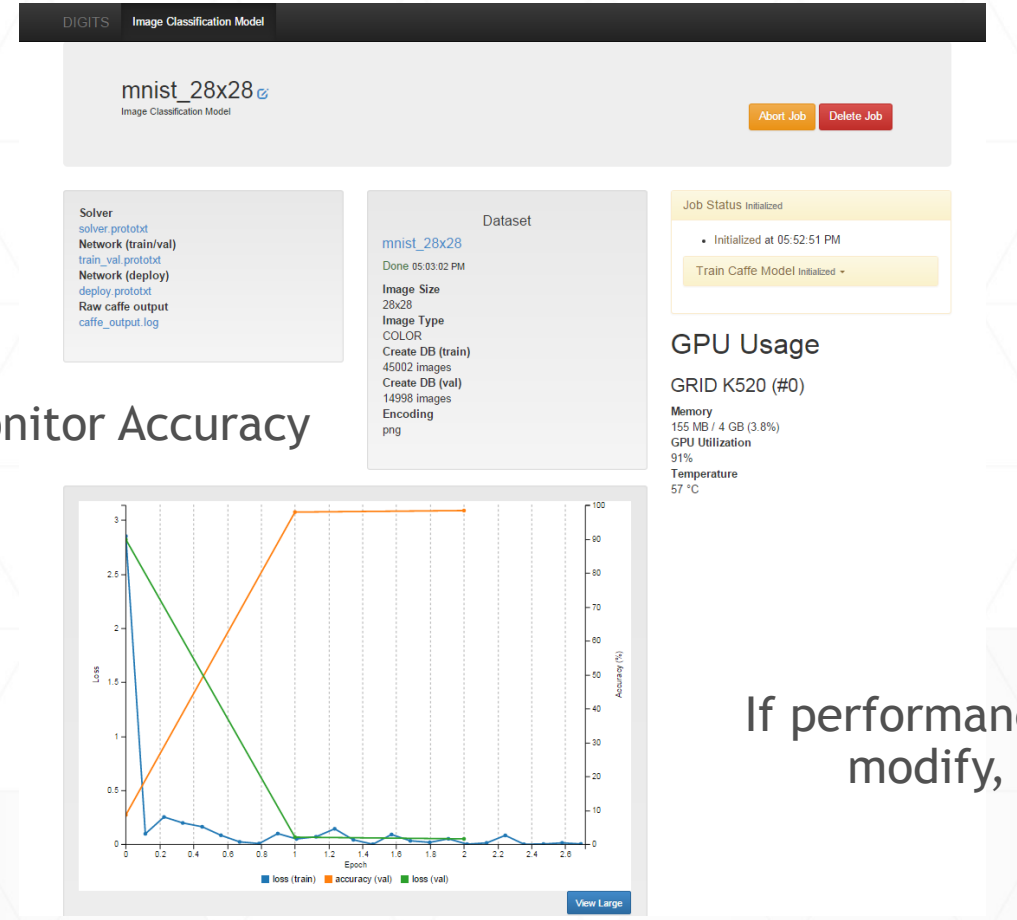

Monitor Accuracy

If performance is poor, abort, modify, and retrain.

OVERFITTING AND UNDERFITTING

How can I use DIGITS to tell me this is happening?

Overfitting

Loss continues to decrease with training data but increases with validation

Validation data helps you identify when/if this occurs!

Underfitting

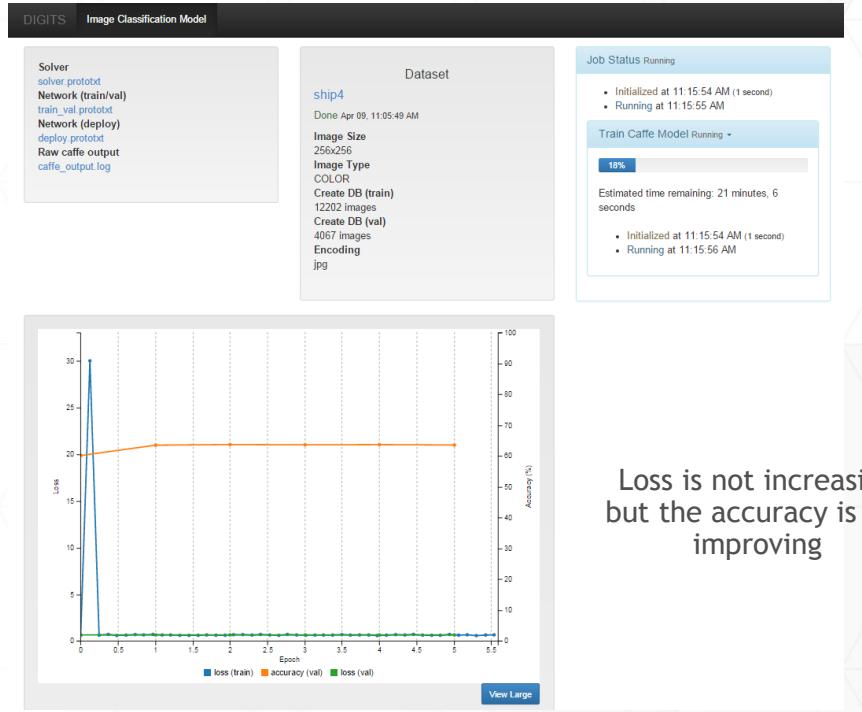

Loss is not increasing but the accuracy is not improving

NVIDIA DIGITS

Single Image Classification

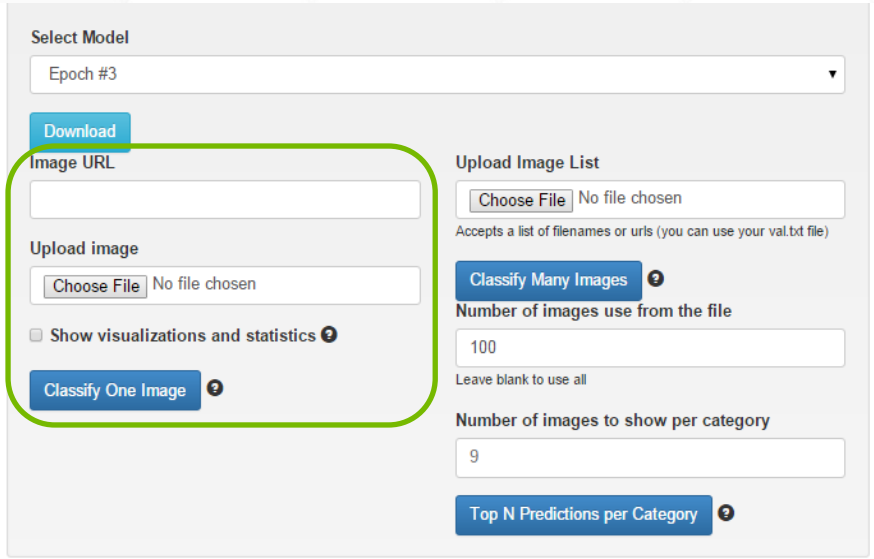

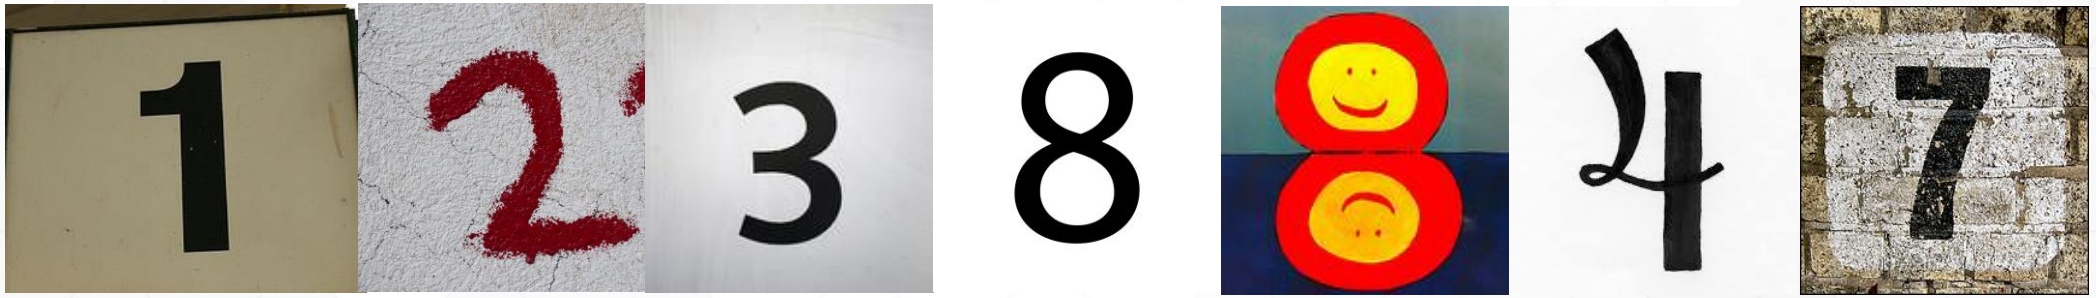

* Lab tasks

SINGLE IMAGE CLASSIFICATION RESULTS

- Network response at each layer will display
- Visualize responses from different inputs

NETWORK CONFIGURATION

Visualizing LeNet Network Responses

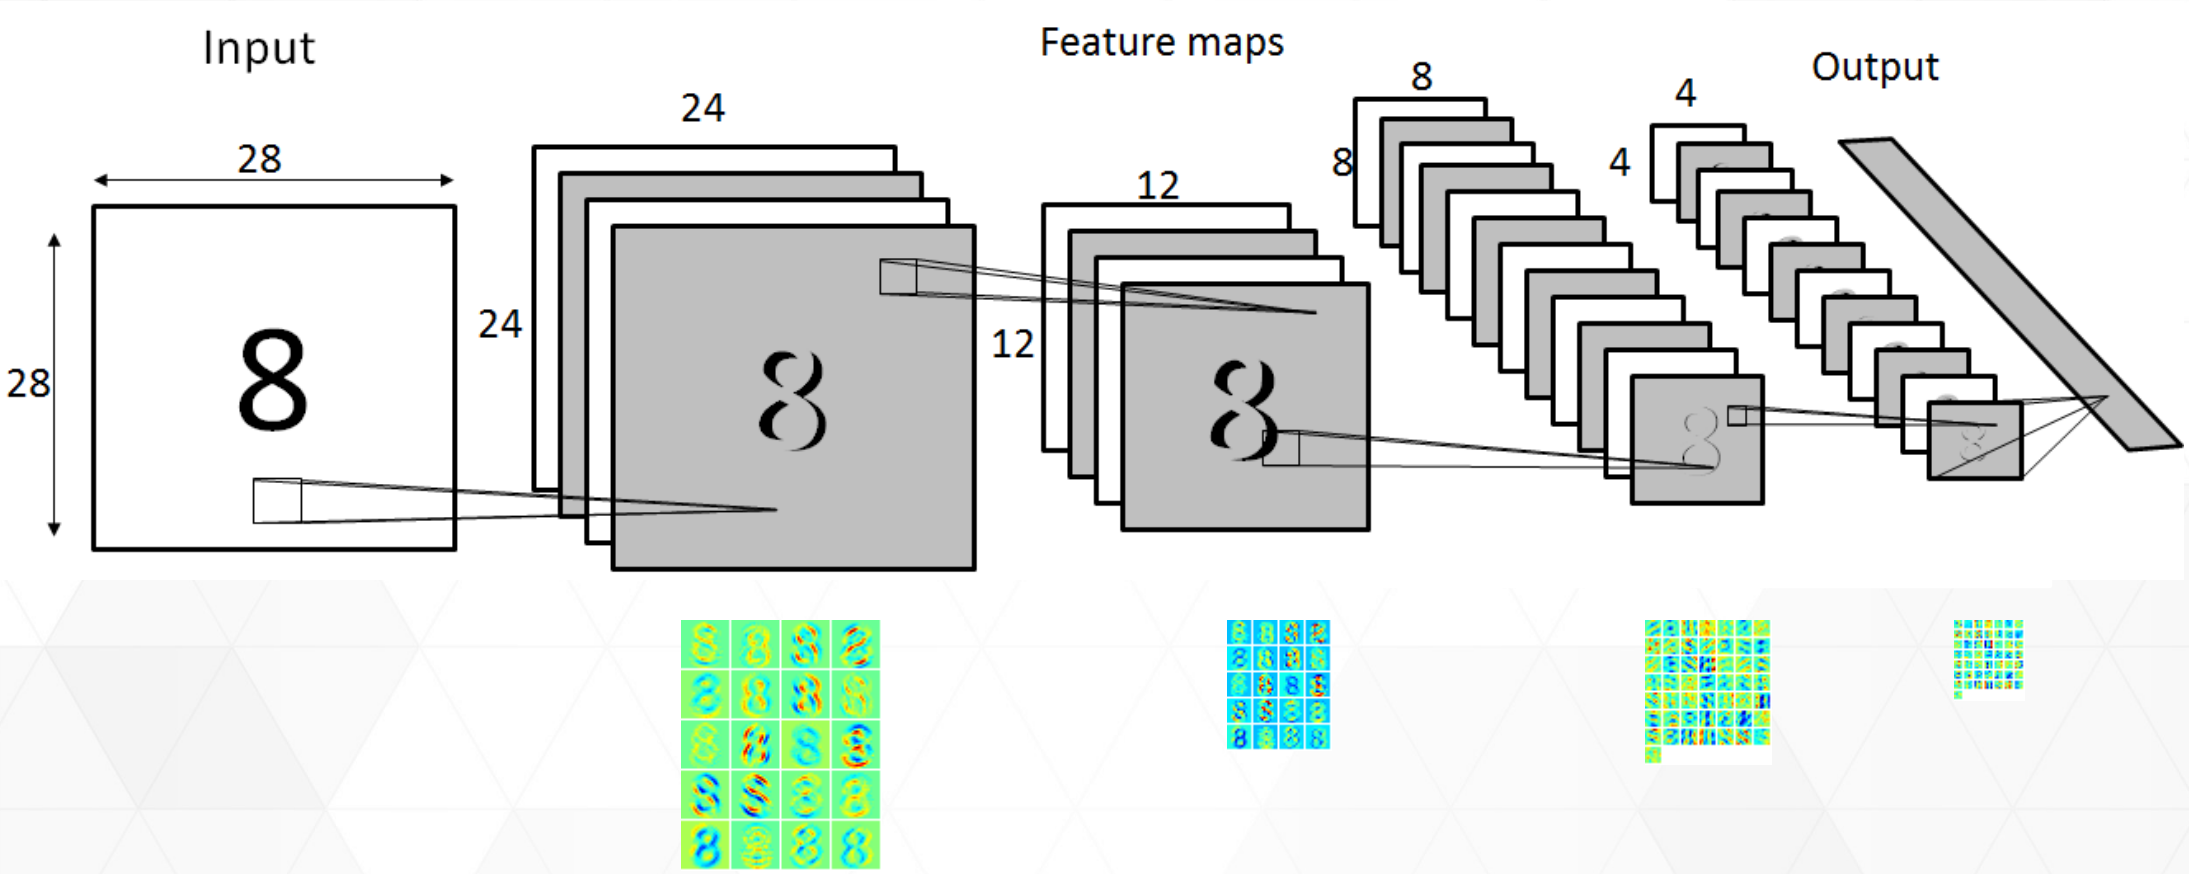

NETWORK CONFIGURATIONS

Modifying a Network

- Modifying a network can improve performance
- There are many parameters - add or remove a layer, pooling, activation function, zero padding, increasing outputs

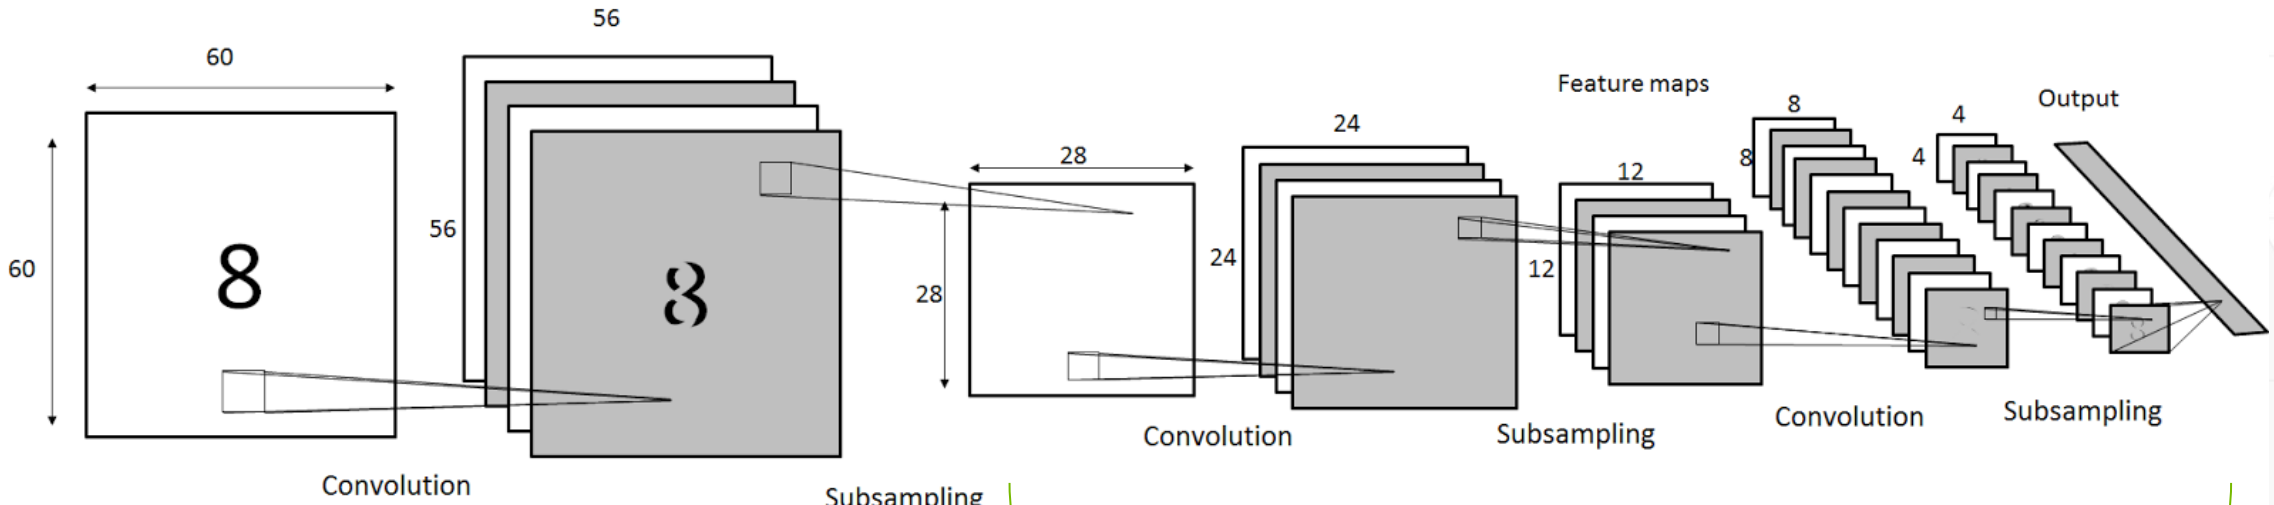

Modification – added a convolutional layer and pooling

LetNet

MODIFYING YOUR NETWORK

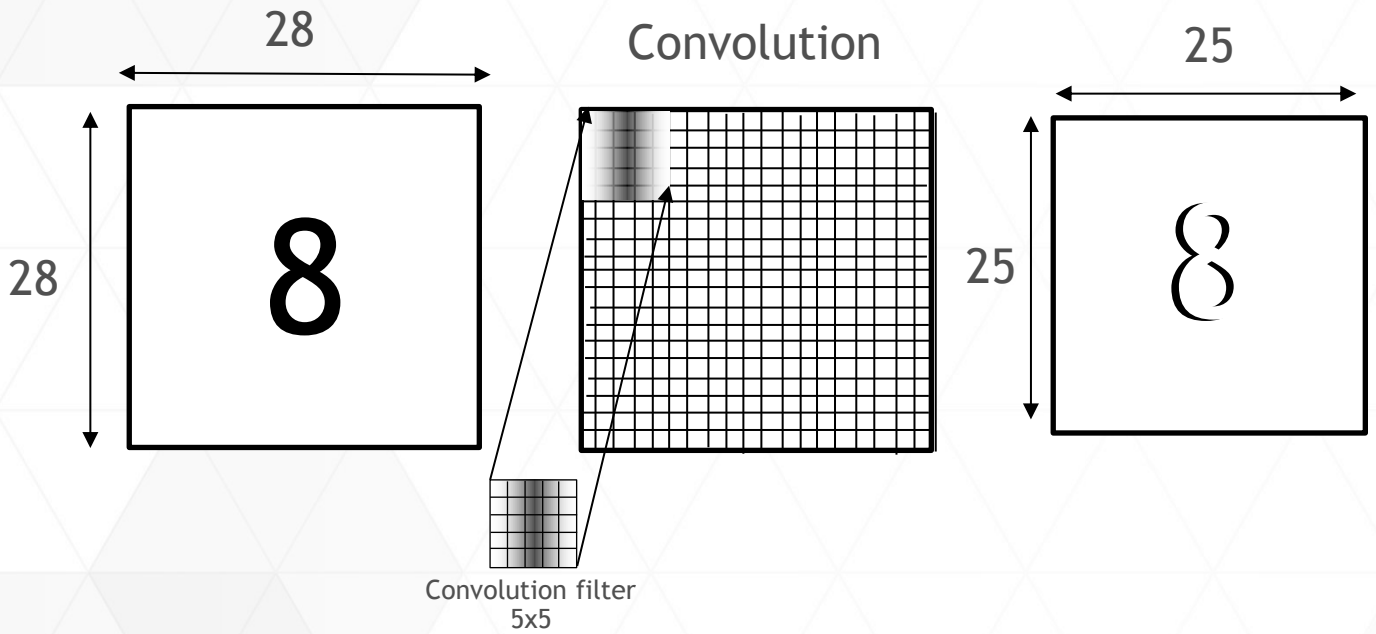

Zero padding, input is reduced by the radius of the kernel
Input[28x28]*filter[5x5]=FeatureMap[25x25]

```
layer {
  name: "conv1"
  type: "Convolution"
  bottom: "scale"
  top: "conv1"
  param {
    lr_mult: 1
  }
  param {
    lr_mult: 2
  }
  convolution_param {
    num_output: 20
    kernel_size: 5
    stride: 1
    weight_filler {
      type: "xavier"
    }
    bias_filler {
      type: "constant"
    }
  }
}
```

MODIFYING YOUR NETWORK

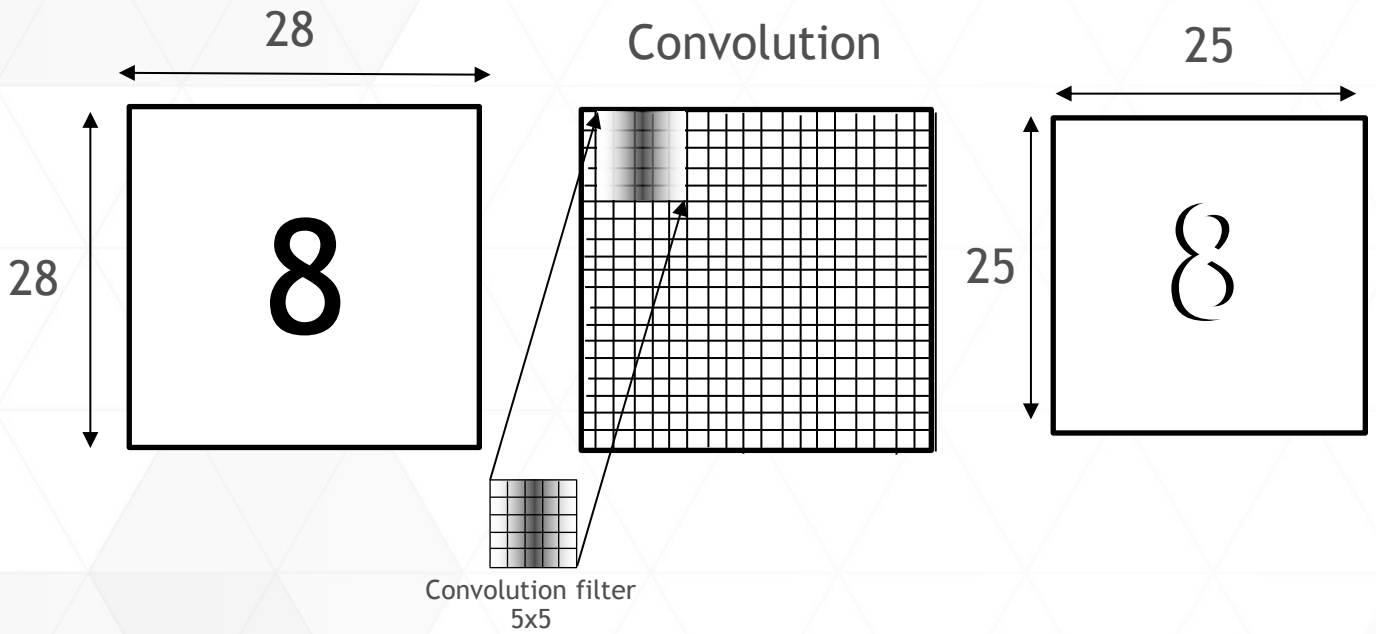

Zero padding, input is reduced by the radius of the kernel
Input[28x28]*filter[5x5]=FeatureMap[25x25]

```
layer {
  name: "conv1"
  type: "Convolution"
  bottom: "scale"
  top: "conv1"
  param {
    lr_mult: 1
  }
  param {
    lr_mult: 2
  }
  convolution_param {
    num_output: 20
    kernel_size: 5
    stride: 1
    weight_filler {
      type: "xavier"
    }
    bias_filler {
      type: "constant"
    }
  }
}
```

MODIFYING YOUR NETWORK

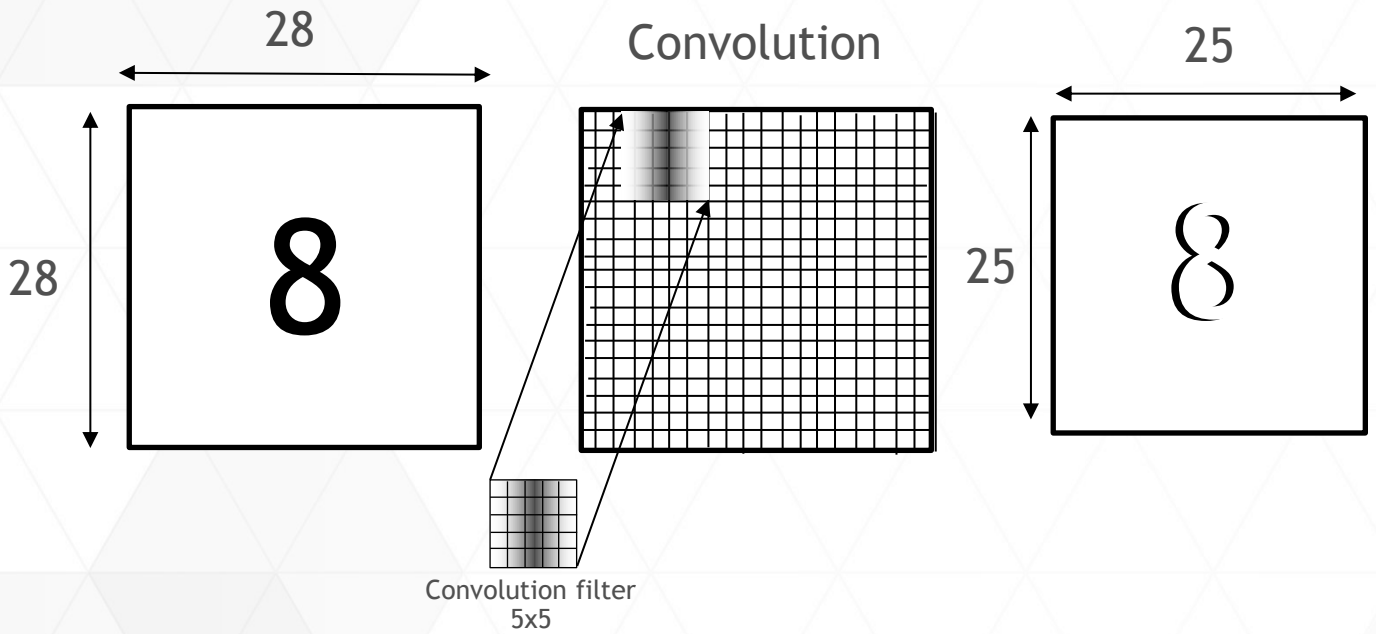

Zero padding, input is reduced by the radius of the kernel
Input[28x28]*filter[5x5]=FeatureMap[25x25]

```
layer {
  name: "conv1"
  type: "Convolution"
  bottom: "scale"
  top: "conv1"
  param {
    lr_mult: 1
  }
  param {
    lr_mult: 2
  }
  convolution_param {
    num_output: 20
    kernel_size: 5
    stride: 1
    weight_filler {
      type: "xavier"
    }
    bias_filler {
      type: "constant"
    }
  }
}
```

MODIFYING YOUR NETWORK

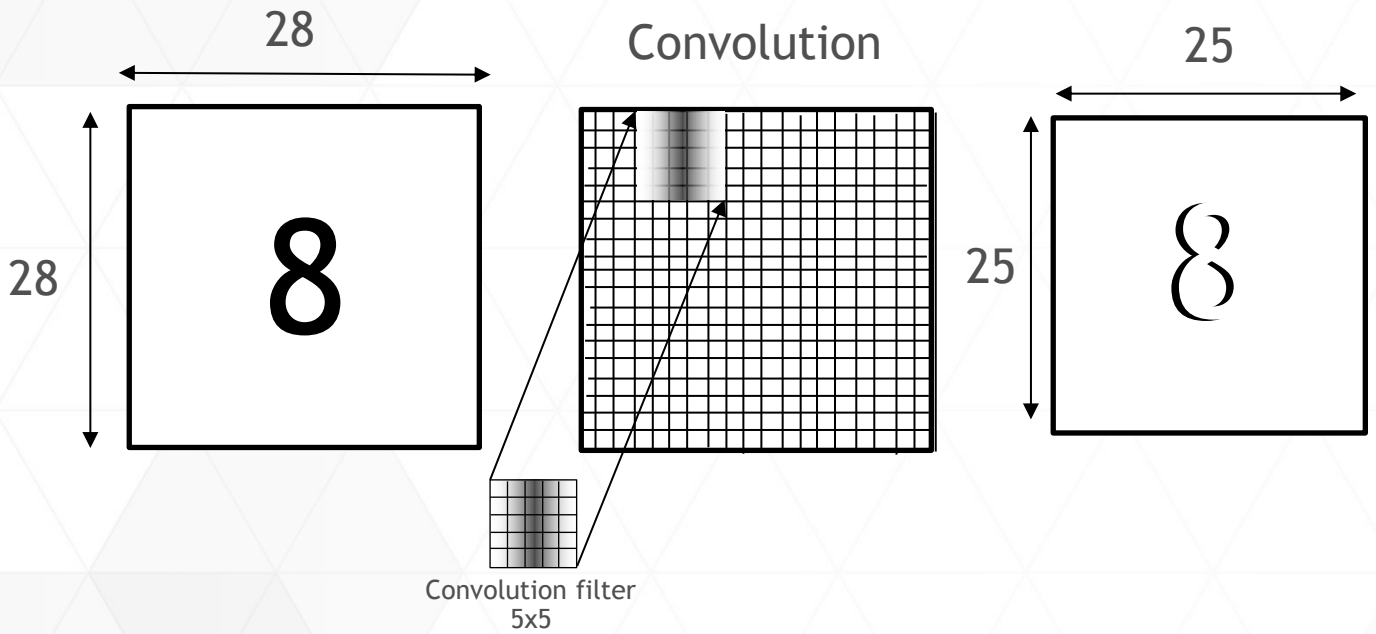

Zero padding, input is reduced by the radius of the kernel
Input[28x28]*filter[5x5]=FeatureMap[25x25]

```
layer {
  name: "conv1"
  type: "Convolution"
  bottom: "scale"
  top: "conv1"
  param {
    lr_mult: 1
  }
  param {
    lr_mult: 2
  }
  convolution_param {
    num_output: 20
    kernel_size: 5
    stride: 1
    weight_filler {
      type: "xavier"
    }
    bias_filler {
      type: "constant"
    }
  }
}
```

NETWORK PARAMETERS

Convolution

```
layer {
 name: "conv0"
 type: "Convolution"
 bottom: "data"
 top: "conv0"
 param {
   lr_mult: 1.0
 }
 param {
   lr_mult: 2.0
 }
 convolution_param {
   num_output: 20
   kernel_size: 5
   stride: 1
   weight_filler {
     type: "xavier"
   }
   weight_filler {
      type: "xavier"
    }
   bias_filler {
      type: "constant"
      value: 0.9
    }
  }
}
```

Pooling/Subsampling

```
layer {
 name: "pool0"
 type: "Pooling"
 bottom: "conv0"
 top: "pool0"
 pooling_param {
   pool: MAX
   kernel_size: 2
   stride: 2
 }
}
```

Activation

```
layer {
 name: "relu0"
 type: "ReLU"
 bottom: "pool0"
 top: "pool0"
}
```

* Lab tasks

NVIDIA DIGITS

Modifying your Network

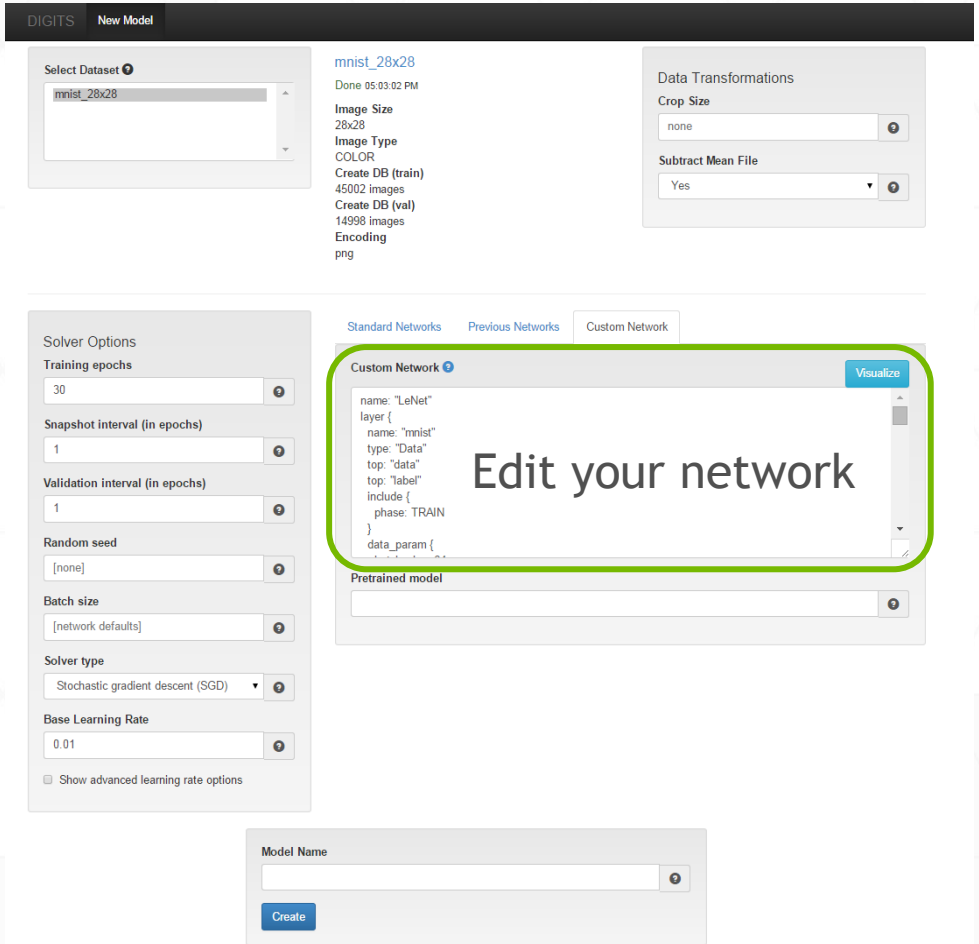

MODIFYING YOUR NETWORK

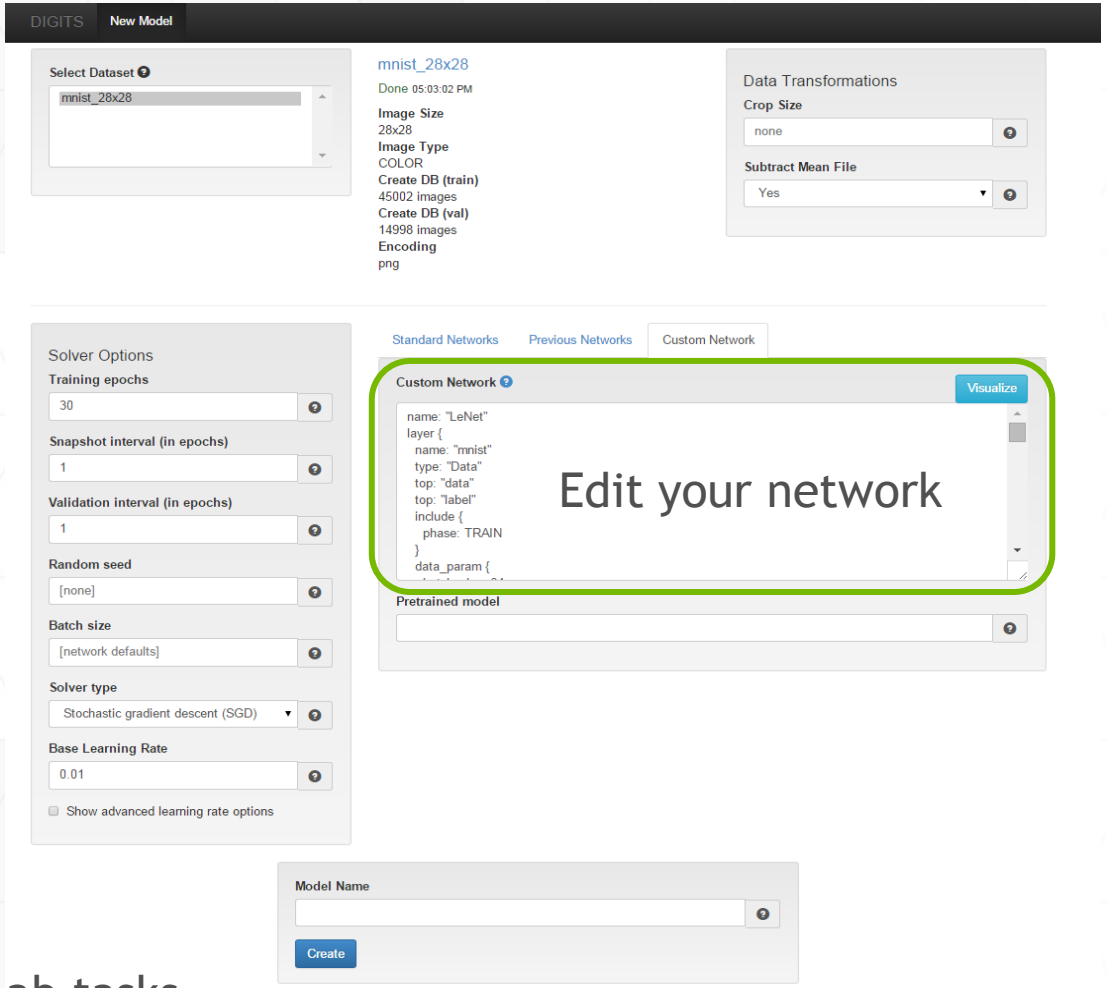


```
layer {
  name: "conv1"
  type: "Convolution"
  bottom: "scale"
  top: "conv1"
  param {
    lr_mult: 1
  }
  param {
    lr_mult: 2
  }
  convolution_param {
    num_output: 20 *
    kernel_size: 5
    stride: 1
    weight_filler {
      type: "xavier"
    }
    bias_filler {
      type: "constant"
    }
  }
}
```

Input and output to the layer

* Lab tasks

MODIFYING YOUR NETWORK

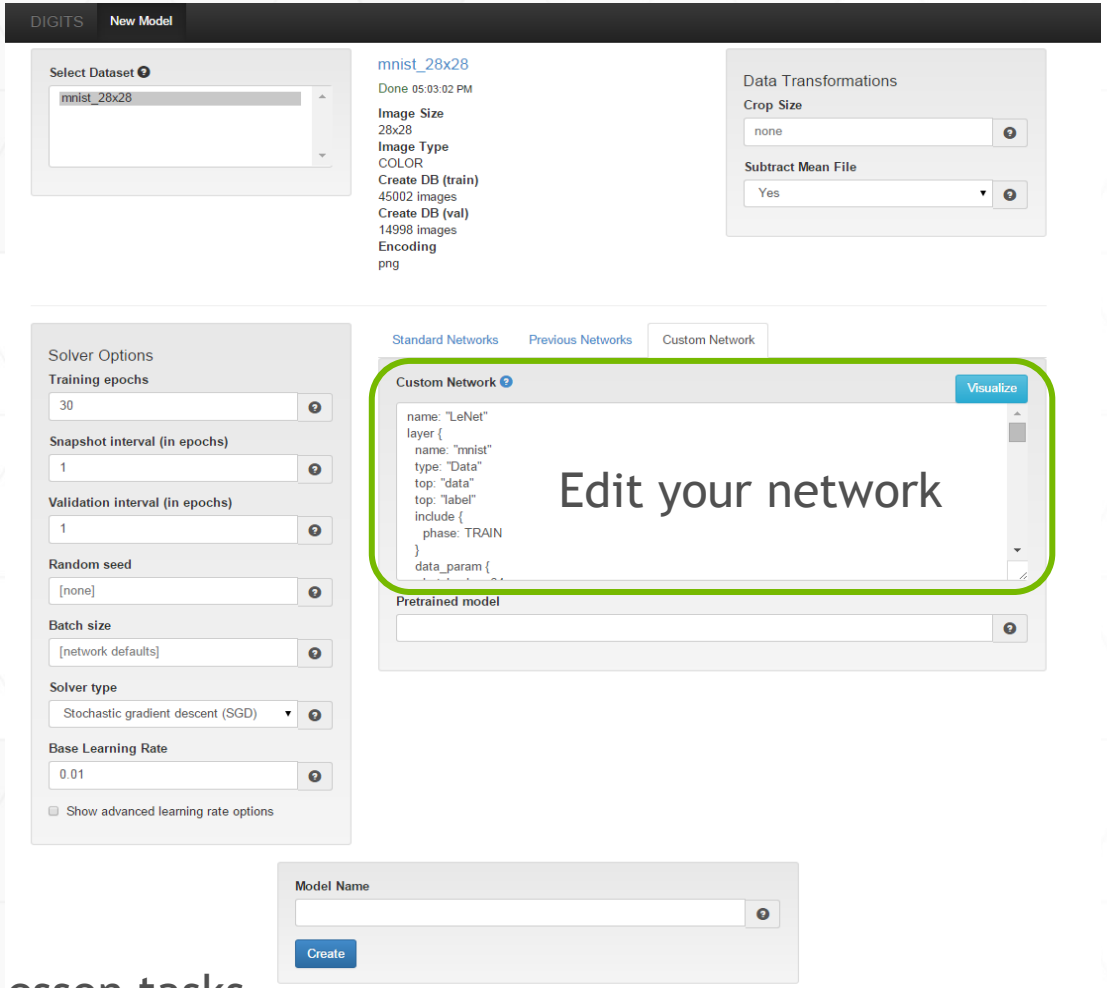

* layer {
  name: "relu1"
  type: "ReLU"
  bottom: "conv1"
  top: "conv1"
  }

Input and output to the layer

* Lesson tasks

MODIFYING YOU NETWORK

Visualize Configuration Changes

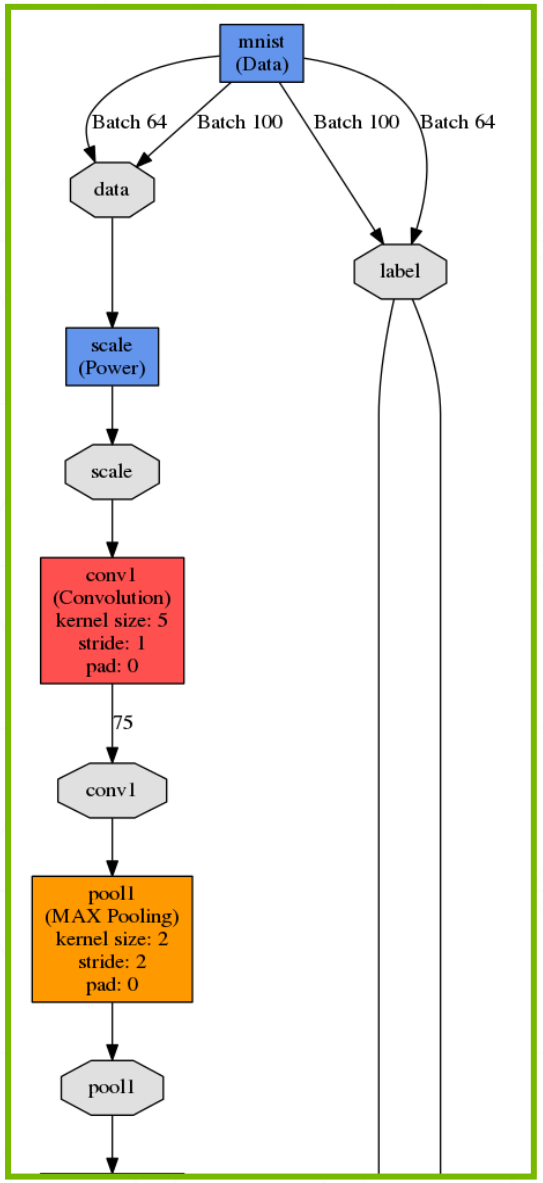

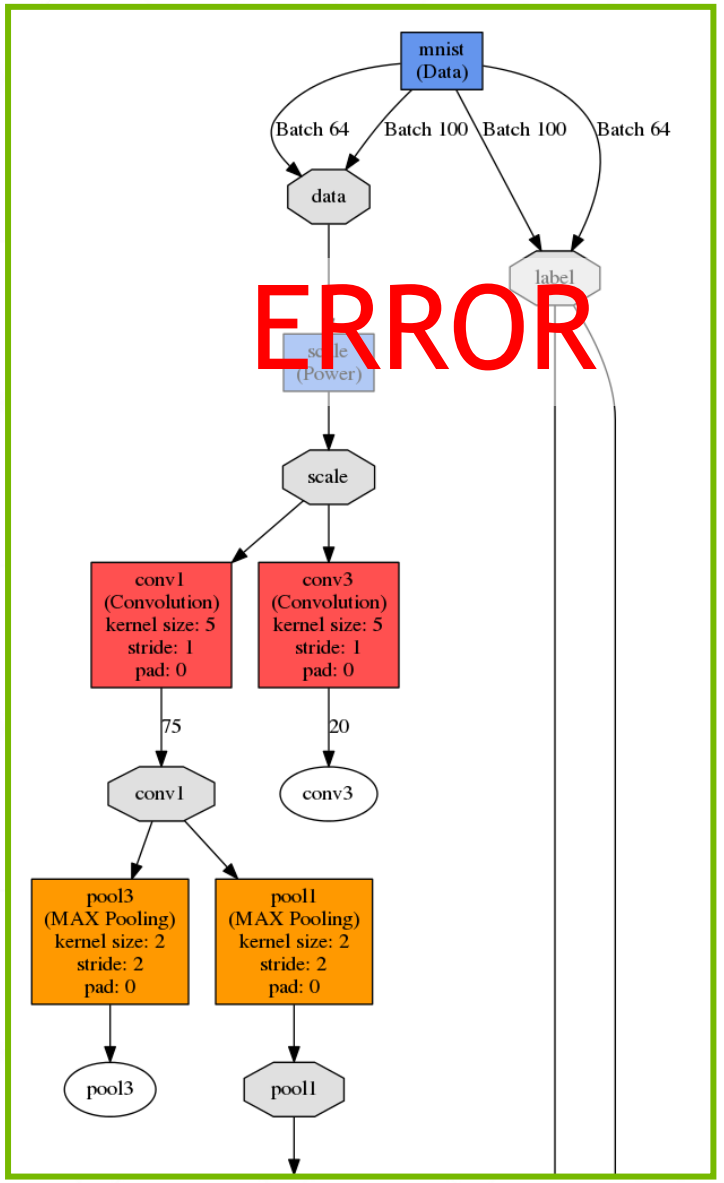

ANOTHER WAY TO IMPROVE PERFORMANCE

Data Augmentation

- Sometimes training data is not a great representation of the field data
  - MNIST data is grayscale, black text with white background

- Will these images to be classified correctly when the network is trained with this digit data?

ANOTHER WAY TO IMPROVE PERFORMANCE

Data Augmentation

- Depending on the deployment scenario, simple modifications can be made to the training data to improve performance

- There are many ways to augment data
  - Rotations, noise, color distortions, stretching, etc. *
- Many ways to modify images – ImageMagick, Python Pillow, OpenCV

* Lesson tasks

USING AN AUGMENTED DATA SET

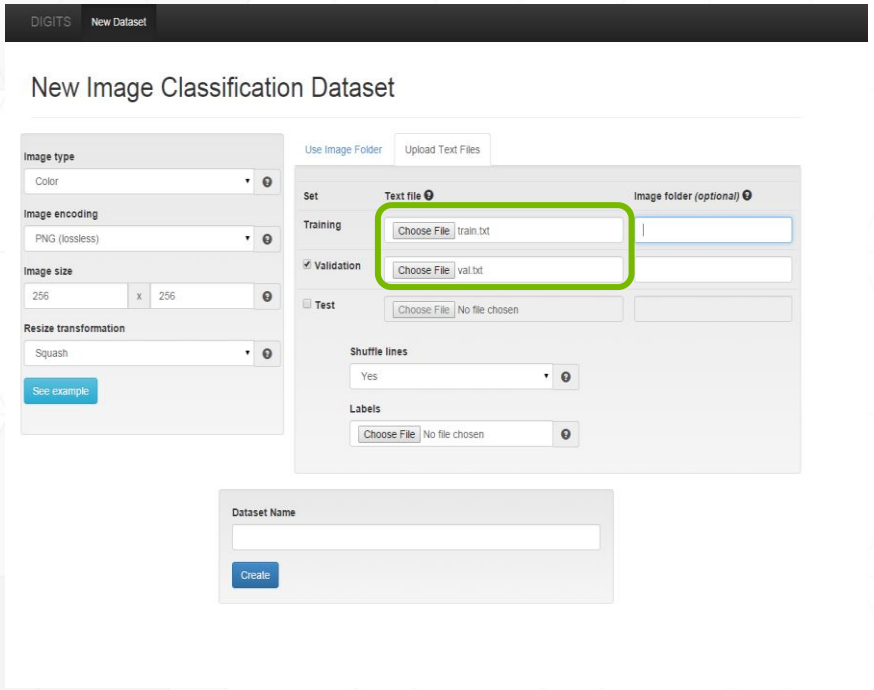

- train.txt

/home/user/train/0/0_1.jpg 0

/home/user/train/0/0_1_invert.jpg 0

/home/user/train/5/5_1.jpg 5

/home/user/train/5/5_1_invert.jpg 5

- val.txt

/home/user/mnist/val/7_1.jpg 7

/home/user/mnist/val/7_1_invert.jpg 7

ANOTHER WAY TO IMPROVE PERFORMANCE

Data Augmentation

- Example augmentation - inverted copies of the input data

- Would a network trained with this data augmented, accurately classify these images?

DEPLOYING YOUR NETWORK

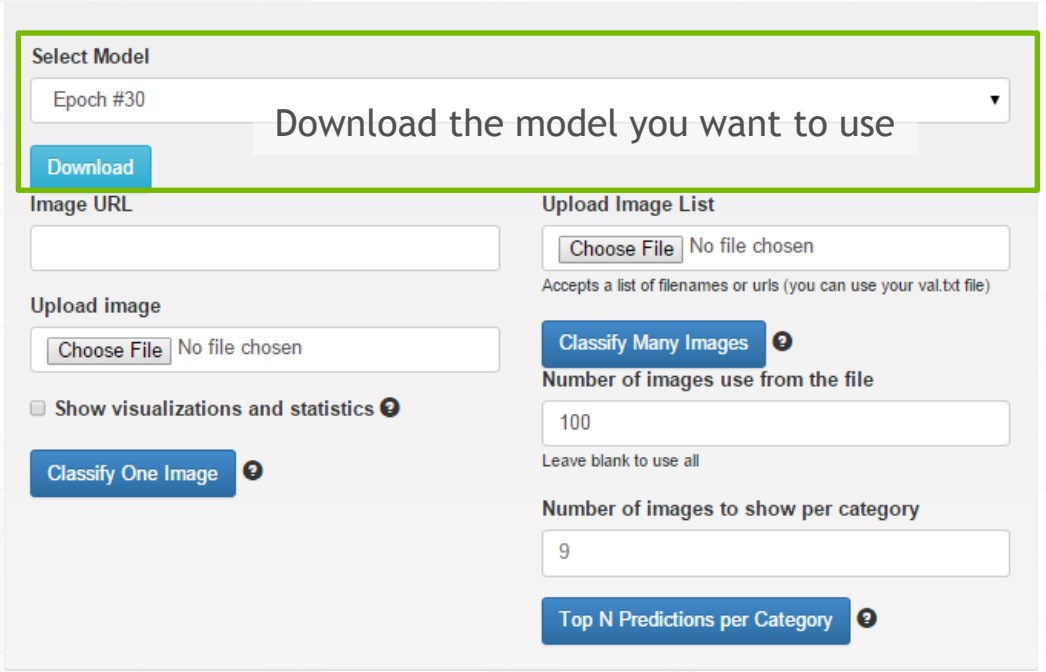

- Deploy in the cloud

- Deploy on a mobile device

Tegra

DEPLOYMENT WITH TEGRA

Rapid Classification Anywhere

- Flexible
- Low Power
- Easy to use
- GPU accelerated

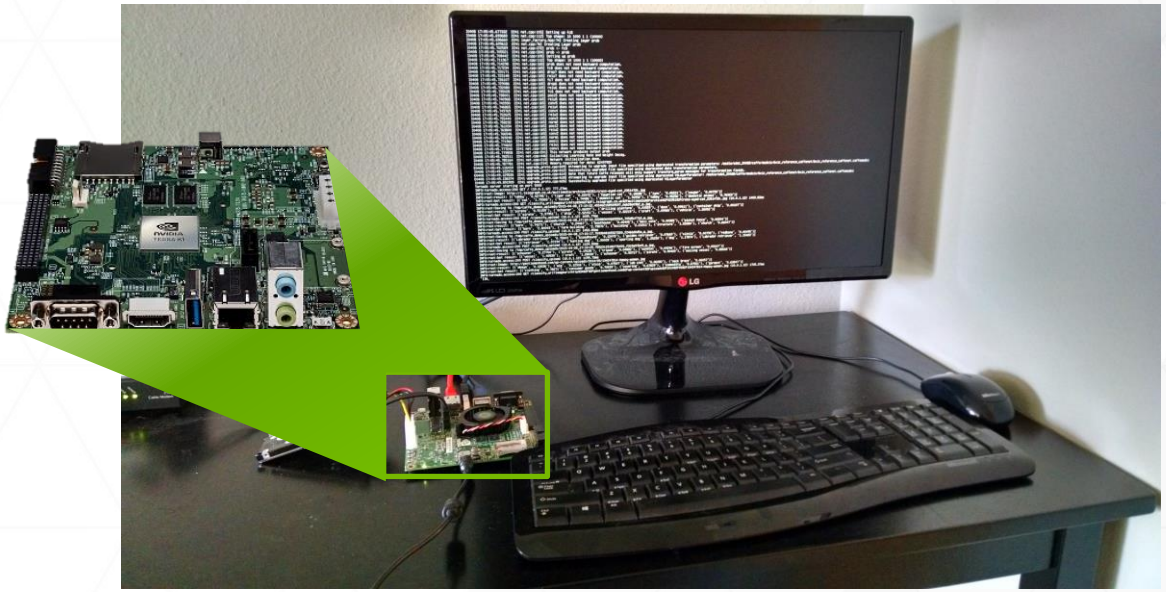

Build Caffe on your portable platform
Download your trained network from DIGITS

NVIDIA DIGITS

Resources

- Where to get DIGITS
  - Easy to use web installer https://developer.nvidia.com/digits
  - github - https://github.com/NVIDIA/DIGITS
    - Remember to install NVIDIA's Caffe branch - https://github.com/NVIDIA/caffe
- User support
  - DIGITS Users Google group - https://groups.google.com/forum/#!forum/digits-users
- For more information on getting started with DIGITS
  - Parallel forall blogs - http://devblogs.nvidia.com/parallelforall/easy-multi-gpu-deep-learning-digits-2/
  - Getting started guide - https://github.com/NVIDIA/DIGITS/blob/master/docs/GettingStarted.md

HANDS-ON LAB

1. Create an account at https://nvidia.qwiklab.com/
2. Go to "Getting Started with DIGITS" lab at
3. Start the lab and enjoy!

- Only requires a supported browser, no NVIDIA GPU necessary!
- Lab is free until end of this Deep Learning Lab series

DEEP LEARNING LAB SERIES SCHEDULE

- All classes start at 9am PT

- 8/5 Class #2 - Getting Started with DIGITS interactive training system for image classification
- 8/12 Office Hours for Class #2

- 8/19 Class #3 - Getting Started with the Caffe Framework
- 8/26 Office Hours for Class #3

- 9/2 Class #4 - Getting Started with the Theano Framework
- 9/9 Office Hours for Class #4

- 9/16 Class #5 - Getting Started with the Torch Framework
- 9/23 Office Hours for Class #5

- More information available at developer.nvidia.com/deep-learning-courses